

TELECOM CUSTOMER CHURN PREDICTION & RETENTION INTELLIGENCE SYSTEM (TelcoPulse AI)

A Machine Learning & Explainable Decision Support System for Proactive Customer Retention

Student Name: Anaganti Sairishikesh

College Name: Engineering & Technology Institute / University

Department: Department of Computer Science and Engineering

Academic Year: Final / Pre-Final Year (2025 - 2026)

Project Domain: Artificial Intelligence & Machine Learning (AI/ML)

Application Deployment: <https://telecomcustomers.streamlit.app/>

TelcoPulse AI • Production-Ready Machine Learning Solution

Table of Contents

S.No	Topic / Chapter Name	Page No.
1	Introduction (Domain, Problem Background, AI/ML Importance, Proposed System)	3
2	Problem Statement (Core Problem, Stakeholders, Existing Limitations, ML Need)	4
3	Objectives (Research & Engineering Milestones)	4
4	Project Scope (Target Users, Application Area, Functional Boundaries)	5
5	Proposed System / Methodology (End-to-End Machine Learning Pipeline)	5
6	Dataset Description (IBM Telco Churn Dataset & Feature Dictionary)	6
7	Data Preprocessing (Cleaning, Imputation, Encoding & Normalization)	7
8	AI/ML Model Selection (Candidate Benchmark & Logistic Regression Selection)	7
9	Model Development (Hyperparameters, Feature Matrices & Development Stack)	8
10	Model Training and Evaluation (Classification Report & Confusion Matrix)	8
11	Results and Analysis (Confusion Matrix, ROC Curve & Log-Odds Feature Importance)	9
12	System Architecture / Workflow (4-Tier Enterprise Decision Support Pipeline)	10
13	Implementation (Software Stack & Deployment Infrastructure)	10
14	User Interface / Application Screenshots (Interactive Web Modules 14.1 to 14.6)	11
15	Challenges and Limitations (Engineering Solutions & Dataset Boundaries)	17
16	Conclusion (Key Project Deliverables & Business Impact)	17
17	Future Scope (Streaming Data, Ensembles & Autonomous CRM Dispatch)	18
18	References (Literature & Scientific Documentation)	18

EXECUTIVE SUMMARY: TelcoPulse AI delivers 80.5% test accuracy and 0.845 ROC-AUC on unseen subscriber telemetry, transforming reactive churn losses into explainable, proactive retention workflows.

1. Introduction

1.1 Brief Introduction to the Project Domain

In the modern telecommunications landscape, subscriber retention has become the primary determinant of long-term corporate viability and market share. As telecom markets reach saturation, acquiring a new customer costs 5 to 7 times more than retaining an existing subscriber. Consequently, telecom operators must transition from reactive damage control to proactive, data-driven customer relationship management.

1.2 Background of the Problem

Customer churn represents the rate at which subscribers cancel their subscriptions, disconnect lines, or switch to competitor networks. Telecom providers historically relied on post-cancellation exit interviews and broad-spectrum marketing discounts. These traditional mechanisms fail because they address customer dissatisfaction only after the cancellation intent is finalized, resulting in billions of dollars in lost annual recurring revenue.

1.3 Importance of AI/ML in the Selected Domain

Artificial Intelligence and Machine Learning (AI/ML) transform retention strategies by detecting subtle, non-linear churn risk patterns hidden within multidimensional customer data. By analyzing demographic details, contracted services, payment behaviors, and billing telemetry, supervised machine learning algorithms can predict customer defection months ahead of time, enabling timely, targeted, and financially optimized retention interventions.

1.4 Brief Overview of the Proposed Solution (TelcoPulse AI)

TelcoPulse AI is an enterprise-grade customer churn intelligence and decision support platform. Rather than acting as an inscrutable black box, TelcoPulse AI provides:

- **Calibrated Risk Scoring:** Real-time calibrated churn risk probability (%) categorized into Safe (<30%), Moderate (30-60%), and Critical (>60%) risk tiers.
- **Local Explainability (XAI):** Mathematical log-odds decomposition identifying the specific hazard factors driving risk and protective factors anchoring loyalty for each subscriber.
- **Financial Exposure Estimation:** Quantifies expected financial loss (\$/year) per customer account based on monthly recurring revenue and churn probability.
- **What-If Strategy Simulator:** An interactive sandbox allowing retention agents to test contract extensions, discounts, and support add-ons to simulate risk reduction before presenting offers.
- **Cohort Batch Queue:** Batch roster upload (.csv) with priority queue sorting and 1-click campaign list export for operational retention workflows.

2. Problem Statement

2.1 What Problem Are We Trying to Solve?

We are addressing the critical problem of voluntary subscriber churn in the telecommunications industry, where customers silently defect to competitor providers without prior warning, inflicting massive revenue erosion and high customer acquisition reinvestment cycles.

2.2 Who Is Affected by the Problem?

- **Customer Care & Retention Specialists:** Struggle with lack of real-time insights during subscriber inbound calls.
- **Marketing & Growth Teams:** Suffer from wasted budget on uniform, untargeted discount campaigns.
- **Executive Leadership & Finance:** Face eroding average revenue per user (ARPU) and customer lifetime value (LTV).
- **End Consumers:** Experience unresolved service friction leading to defection.

2.3 What Are the Limitations of the Existing Approach?

- **1. Reactive Detection:** Churn is identified only after the customer requests service termination or stops paying bills.
- **2. Blanket Discounting Inefficiencies:** Standard retention campaigns distribute blanket promotions to entire subscriber bases, eroding profit margins on loyal customers who would not have left.
- **3. Black-Box Opacity:** Conventional churn models output raw probability numbers without explaining why the customer is leaving, reducing agent trust.
- **4. Lack of Simulation Capabilities:** Retention agents have no dynamic tools to evaluate the potential impact of retention packages before making counter-offers.

2.4 Why Is an AI/ML-Based Solution Required?

Subscriber attrition involves complex, non-linear interactions between tenure, contract type, internet technology, payment methods, and bundled add-on services. Human intuition and rule-based heuristics cannot reliably weigh these 20+ interacting variables simultaneously. An AI/ML-based classifier trained on historical subscriber data learns the multidimensional decision boundaries necessary to accurately predict churn probability and uncover actionable retention levers.

3. Objectives

- **1. Develop AI/ML Solution:** Develop an end-to-end AI/ML-based decision support system to accurately classify customer churn risk.
- **2. Analyze Available Data:** Perform thorough Exploratory Data Analysis (EDA) on telecom subscriber records to uncover underlying behavioral patterns and risk factors.
- **3. Build & Train Suitable Model:** Train, optimize, and benchmark multiple supervised classification models to establish the most effective and interpretable architecture.
- **4. Evaluate Model Performance:** Rigorously validate model performance using Accuracy, Precision, Recall, F1-Score, ROC-AUC, and Confusion Matrix metrics.
- **5. Provide Actionable Predictions & Insights:** Deliver calibrated predictions, localized explainability (log-odds attribution), financial risk exposure, and interactive What-If simulation via a production web dashboard.

4. Project Scope

The scope and operational boundaries of the TelcoPulse AI project comprise:

- **Target Users:** Frontline Telecom Customer Support Representatives, Dedicated Retention Agents, Marketing Campaign Managers, and Executive Operations.
- **Application Area:** Telecommunication service providers managing broadband (DSL, Fiber optic), fixed-line telephony, and bundled digital value-added services.
- **Features Covered:** Single Customer Risk Diagnosis, What-If Strategy Simulator, Batch File Scoring (.csv), Executive Cohort Analytics, and Model Transparency Diagnostics.
- **Expected Outputs:** Calibrated Churn Probability (%), Risk Tier Categorization, Annual Financial Loss Estimation (\$), Top Risk Factor Drivers, and Prescriptive Playbook Actions.
- **Current Project Limitations:** The project operates on static tabular snapshot data and does not ingest live continuous streaming network telemetry logs.

5. Proposed System / Methodology

The overall end-to-end Machine Learning lifecycle is structured according to the following standardized sequential pipeline:

Data Collection → Data Preprocessing → Feature Engineering → Model Selection → Model Training → Model Evaluation → Prediction & Dashboard Output

Methodology Phase	Key Activities & Technical Outcomes
Phase 1: Data Collection & Ingestion	Acquire the benchmark IBM Telco Customer Churn dataset comprising 7,043 customer accounts and 21 attributes.
Phase 2: Data Preprocessing	Clean missing whitespace anomalies in TotalCharges, verify duplicate records, and audit outlier distributions.
Phase 3: Feature Engineering	Apply OneHotEncoder with drop-first logic to categorical variables and StandardScaler to normalize continuous numerical features.
Phase 4: Model Selection & Training	Train and benchmark multiple candidate classifiers (Logistic Regression, Random Forest, Decision Tree, SVM, KNN) using stratified 80/20 train-test splits.
Phase 5: Model Evaluation	Evaluate models across Accuracy, Precision, Recall, F1-Score, ROC-AUC, and Confusion Matrix to select the optimal model.
Phase 6: Explainability Integration	Extract beta coefficients and log-odds contributions to build localized risk attribution bars for every individual inference.
Phase 7: Web Dashboard & Deployment	Develop an interactive decision support dashboard using Streamlit and deploy the application live to Streamlit Cloud.

6. Dataset Description

The project utilizes the standardized IBM Telecommunications Customer Churn dataset. The dataset provides comprehensive cross-sectional records of 7,043 unique residential telecom subscribers across 21 attributes covering demographics, subscribed services, account tenure, payment methods, and historical churn status.

DATASET SPECIFICATIONS: Source: IBM Analytics / Kaggle Telco Churn Repository | Total Records: 7,043 | Attributes: 21 | Target: Churn (Yes: 1,869 [26.54%], No: 5,174 [73.46%])

6.1 Sample Dataset Records (First 5 Rows)

customerID	gender	tenure	Contract	PaymentMethod	Monthly	Total	Churn
7590-VHVEG	Female	1 mo	Month-to-month	Electronic check	\$29.85	\$29.85	No (0)
5575-GNVDE	Male	34 mo	One year	Mailed check	\$56.95	\$1889.50	No (0)
3668-QPYBK	Male	2 mo	Month-to-month	Mailed check	\$53.85	\$108.15	Yes (1)
7795-CFOCW	Male	45 mo	One year	Bank transfer	\$42.30	\$1840.75	No (0)
9237-HQITU	Female	2 mo	Month-to-month	Electronic check	\$70.70	\$151.65	Yes (1)

6.2 Feature Attribute Dictionary

Feature / Group Name	Data Type	Description / Value Domain
customerID	Categorical (ID)	Unique identifier assigned to each subscriber account
gender / SeniorCitizen	Binary	Gender (Male/Female) and Senior status (1 = 65+, 0 = Under 65)
Partner / Dependents	Binary	Whether subscriber has a spouse/partner or live-in dependents
tenure	Numerical (Int)	Number of active subscription months with company (0 - 72)
PhoneService / MultipleLines	Categorical	Fixed-line telephone service (Yes/No) and multi-line availability
InternetService	Categorical (3 levels)	Broadband technology: DSL, Fiber optic, or No internet
Value-Added Services (6)	Categorical (3 levels)	Online Security, Backup, Device Protection, Tech Support, Streaming
Contract	Categorical (3 levels)	Billing commitment: Month-to-month, One year, or Two year
PaperlessBilling / Payment	Categorical	Paperless status and Payment Method (Electronic check, Card, Bank)
MonthlyCharges / TotalCharges	Numerical (Float)	Current monthly bill (\$18.25 - \$118.75) and cumulative billing
Churn (Target)	Binary (0 / 1)	Ground truth label: whether subscriber churned (Yes = 1, No = 0)

7. Data Preprocessing

Data preprocessing ensures raw tabular telemetry is transformed into clean, mathematically robust feature matrices suitable for machine learning algorithms:

- **1. Handling Missing Values:** The TotalCharges column contained 11 whitespace strings (' ') corresponding to new customers (tenure = 0). These were converted to numeric float and imputed to \$0.00 without row deletion.
- **2. Duplicate Verification:** The entire dataset was checked for duplicate records across all 21 columns; zero duplicate rows were found.
- **3. Categorical Encoding:** Categorical variables (16 features) were converted using OneHotEncoder with drop='first', producing 41 binary indicator columns to prevent multicollinear dummy traps.
- **4. Feature Standardization:** Continuous numerical features and encoded columns were normalized using StandardScaler to achieve zero mean and unit variance.
- **5. Stratified Partitioning:** Stratified 80/20 train/test split (5,634 train / 1,409 test) preserved the 73.5% vs. 26.5% class distribution across both partitions.

8. AI/ML Model Selection

Five candidate supervised classification algorithms were implemented, tuned, and evaluated on the standardized test set:

Model	Accuracy	Precision	Recall	F1	AUC	Interpretability	Latency
Logistic Regression (Selected)	80.48%	65.2%	55.8%	0.601	0.845	High (Direct Log-Odds)	< 5 ms
Random Forest Classifier	79.20%	62.8%	49.4%	0.553	0.828	Moderate (Tree Ensemble)	35 ms
Decision Tree Classifier	72.82%	48.9%	50.5%	0.497	0.657	High (Single Tree)	< 5 ms
Support Vector Machine (RBF)	79.84%	64.1%	51.2%	0.569	0.812	Low (Black Box)	60 ms
K-Nearest Neighbors (KNN)	76.44%	55.3%	52.7%	0.540	0.781	Low (Instance Based)	25 ms

8.1 Selection Rationale for Production Deployment

- **1. Superior Discriminative Power:** Logistic Regression achieved top test accuracy (80.48%) and the highest Area Under ROC (0.845 AUC).
- **2. Sub-5ms Operational Latency:** Under 5ms inference latency ensures instantaneous real-time responsiveness during live customer service calls.
- **3. Deterministic Explainability:** Direct extraction of learned beta log-odds allows instant decomposition into transparent risk drivers for frontline agents.

9. Model Development

The model development phase encompasses feature pipeline engineering, hyperparameter optimization, and development tooling configuration:

- **Feature Pipeline:** All 20 predictor attributes were transformed into a standardized 41-dimensional numeric feature vector.
- **Hyperparameter Setup:** Penalty = L2 (Ridge regularization), Inverse Regularization (C) = 1.0, Optimization Solver = lbfgs, Max Iterations = 1000, Random State = 42 for deterministic reproducibility.
- **Technology Stack:** Python 3.10+, Scikit-learn 1.3+, Pandas, NumPy, Streamlit, Joblib, Plotly, Matplotlib.

10. Model Training and Evaluation

The training and evaluation lifecycle followed a disciplined cross-validation structure:

Training Data (5,634 rows) → Model Fitting (lbfgs solver) → Validation & Test Data (1,409 rows) → Calibrated Predictions → Metric Evaluation

Class / Metric	Precision	Recall	F1-Score	Support Count
Retained (Class 0: No Churn)	0.85 (85%)	0.90 (90%)	0.87 (87%)	1,035
Churned (Class 1: Churn)	0.66 (66%)	0.56 (56%)	0.60 (60%)	374
Overall Accuracy	-	-	0.805 (80.5%)	1,409
Macro Average	0.76 (76%)	0.73 (73%)	0.74 (74%)	1,409
Weighted Average	0.80 (80%)	0.80 (80%)	0.80 (80%)	1,409

10.1 Confusion Matrix Breakdown

Ground Truth	Predicted: Retained (No)	Predicted: Churned (Yes)
Actual Retained (No Churn)	930 (True Negatives - TN)	105 (False Positives - FP)
Actual Churned (Yes Churn)	165 (False Negatives - FN)	209 (True Positives - TP)

11. Results and Analysis

Model evaluation and exploratory attribution yielded definitive empirical findings regarding subscriber churn dynamics:

11.1 Feature Importance & Log-Odds Risk Attribution




Figure 1: Key Churn Risk Drivers (+) vs Retention Anchors (-)

11.2 Key Analytical Observations

- **Contract Sensitivity:** Month-to-month contracts (+1.28) and electronic check payments (+0.42) are the highest risk indicators. Migrating customers to 1- or 2-year contracts slashes churn by over 70%.
- **Service Bundling:** Fiber optic customers (+0.84) churn at 41.9% due to higher billing friction (\$70-\$100/mo), whereas bundling Tech Support (-0.58) and Online Security (-0.46) reduces churn to 14.6%.
- **Tenure Compounding:** Subscribers who stay beyond 24 months (-0.62) demonstrate over 88% retention, highlighting the critical importance of first-year onboarding.
- **Model Robustness:** Confusion matrix and ROC analysis (AUC = 0.845) confirm robust discrimination between loyal subscribers and defectors across varied thresholds.

12. System Architecture / Workflow

The TelcoPulse AI system is architected as an end-to-end reactive decision support pipeline:

User / Retention Agent --> Single Form / Batch CSV --> Preprocessing & Encoding --> Logistic Regression ML Engine --> Risk Probability & Log-Odds --> Dashboard & Prescriptive Playbook

Architecture Tier	Component Role & Responsibilities
Tier 1: Presentation Tier	Streamlit UI providing Single Customer Diagnosis, What-If Strategy Simulator, Batch File Scoring, Cohort Insights, and Diagnostics.
Tier 2: Inference & Attribution Tier	Logistic Regression Engine coupled with real-time log-odds decomposition and financial risk calculations.
Tier 3: Preprocessing & Serialization	StandardScaler and OneHotEncoder pipelines loaded from serialized joblib artifacts.
Tier 4: Data & Export Tier	Telco Customer CSV store, Batch Scoring Priority Queues, and 1-click campaign CSV export.

13. Implementation

The technical implementation details of the software stack are structured as follows:

- **Programming Language:** Python 3.10+ (Primary computational and web backend language).
- **Frameworks:** Streamlit 1.30+ (Reactive web UI framework) and Scikit-learn 1.3+ (Machine learning framework).
- **Core Libraries:** Pandas, NumPy, Plotly Express, Matplotlib, Joblib, Python-Docx.
- **Deployment Platform:** Streamlit Community Cloud (<https://telecomcustomers.streamlit.app/>).

14. User Interface: 14.1 Single Customer Risk Diagnosis & Speedometer Gauge

Features an interactive risk speedometer (Safe <30%, Moderate 30-60%, Critical >60%), annual financial exposure calculation (\$/year), local log-odds driver attribution waterfall, and AI-recommended prescriptive retention playbooks.

Figure 2: Single Customer Risk Diagnosis & Speedometer Gauge

Key Feature / UI Component	Operational Purpose & Business Value
Interactive Speedometer	Visualizes real-time calibrated churn probability and risk tier classification.
Annual Financial Loss	Calculates revenue at risk (\$/year) based on monthly bill and defection likelihood.
Top Risk Factors Waterfall	Ranks top drivers accelerating churn (e.g. Month-to-Month contract, Fiber optic).
Actionable Playbook	Suggests targeted retention incentives based on identified account vulnerabilities.

14. User Interface: 14.2 GenAI Customer Retention Copilot Interface

Features prompt orchestration, multi-angle stochastic reasoning, Creativity/Temperature slider (0.2–1.0), and 4 generated retention assets: personalized win-back emails, frontline agent call scripts, concession hierarchy, and CRM JSON payload.



TelcoPulse AI

Real-time Retention & Churn Intelligence

Navigation

- Single Customer Diagnostics
- GenAI Retention Copilot
- Virtual AI Retention Sandbox
- Batch Scoring & Priority Queue
- Executive Cohort Insights
- Model Diagnostics

Pre-Tip: Use the GenAI Copilot to automatically generate personalized win-back campaigns and retention scripts.

Deploy

GenAI Customer Retention Copilot

Harness Generative AI to instantly generate personalized win-back emails, frontline customer service negotiation scripts, and CRM-ready retention campaigns tailored to each subscriber's churn risk profile.

1. Customer Context & Risk Telemetry

Customer Name / Account ID

Sarah Jenkins (ID: #7590-VHVEG)

Predictive ML Risk Tier

Critical Risk (>60% Churn Probability)

Customer Tenure (Months)

4

Contract Type

Month-to-month

Monthly Charges (Current)

94.50

Selected Service Type

Fiber optic

Identified Primary Churn Drivers (Based on ML Model)

High Monthly Bill (\$90+) × Month-to-Month Contract Flexibility × No

2. GenAI Campaign Configuration

Campaign Objective

Emergency Win-Back & Churn Reversal Offer

AI Voice & Tone Persona

Empathetic, Warm & Customer-Centric

Primary Communication Channel

Personalized Email Campaign

SMS / Push Notification Alert

Call Center Dialogue Script

Executive Retention Brief

Creativity & Temperature (Variability Level)

0.55

View Structured Prompt Template Sent to LLM

Generate AI Retention Campaign & Playbook

Re-Think & Generate New Angle

Figure 3: GenAI Customer Retention Copilot Interface

Key Feature / UI Component	Operational Purpose & Business Value
Prompt Sidebar	Ingests customer risk tier, tenure, charges, and campaign objectives.
Creativity Slider	Adjusts LLM temperature between deterministic (0.2) and creative (1.0).
Generated Copy Tabs	Provides instant empathetic email copy and frontline call objection rebuttals.
CRM Integration	Outputs validated JSON payload for instant enterprise CRM synchronization.

14. User Interface: 14.3 What-If Retention Strategy Simulator

Enables customer service agents to test contract extensions, technical support bundles, and billing discounts to simulate risk reduction in real time before pitching counter-offers.




Figure 4: What-If Retention Strategy Simulator

Key Feature / UI Component	Operational Purpose & Business Value
Dynamic Contract Toggle	Simulates transitioning a month-to-month subscriber to a 1-year or 2-year plan.
Service Bundle Add-ons	Tests the protective impact of adding Tech Support or Online Security packages.
Instant Risk Delta	Computes immediate percentage drop in churn probability (e.g. 74% -> 22%).
Margin Guardrails	Ensures proposed discounts preserve corporate profitability thresholds.

14. User Interface: 14.4 Batch Cohort Scoring & Prioritized Action Queue

Scores entire subscriber rosters (.csv), ranks accounts by expected annual financial loss, and generates downloadable prioritized campaign lists for targeted marketing outreach.




Figure 5: Batch Cohort Scoring & Prioritized Action Queue

Key Feature / UI Component	Operational Purpose & Business Value
CSV File Uploader	Accepts enterprise billing exports with automated column validation.
Risk Distribution Graph	Visualizes breakdown of Safe, Moderate, and Critical accounts in the roster.
Prioritized Queue	Ranks accounts by expected annual revenue loss to maximize retention ROI.
Exportable Outreach List	Enables 1-click download of filtered subscriber segments for campaign dispatch.

14. User Interface: 14.5 Executive Cohort Analytics & Macro EDA

Presents macro-level retention insights across contract types, payment methods, monthly charges distribution, and tech support adoption.




Figure 6: Executive Cohort Analytics & Macro EDA

Key Feature / UI Component	Operational Purpose & Business Value
Contract Comparison	Shows 42.7% month-to-month churn vs. under 3% on two-year agreements.
Payment Method Breakdown	Reveals electronic check churn (45.3%) is 3x higher than credit card users.
Pricing Sensitivity EDA	Analyzes churn concentration between \$70 and \$100 monthly charge bands.
Tenure Hazard Curves	Illustrates steep risk reduction once subscribers surpass 12 months tenure.

14. User Interface: 14.6 Model Diagnostics & Global Feature Importance

Provides complete transparency into pipeline architecture, preprocessing parameters, and global beta coefficients.




Figure 7: Model Transparency & Global Feature Importance

Key Feature / UI Component	Operational Purpose & Business Value
Pipeline Transparency	Displays full preprocessing transformations and regularization parameters.
Global Coefficients	Visualizes all 41 feature weights learned by the production classifier.
Threshold Tuning	Allows operators to adjust decision thresholds based on precision-recall trade-offs.
Data Audit Metrics	Verifies zero missing value residuals and normalized distribution stats.

15. Challenges and Limitations

15.1 Challenges Encountered

- **Data Quality & Missing Values:** Handling whitespace anomalies in TotalCharges and ensuring clean data conversions without record loss.
- **Class Imbalance:** The dataset contains an inherent 73.5% vs. 26.5% non-churn vs. churn imbalance, resolved using stratified sampling and threshold calibration.
- **Model Selection Trade-off:** Selecting an architecture that balances high predictive accuracy with immediate log-odds explainability for non-technical agents.
- **High Dimensionality & Collinearity:** Transforming 16 categorical features created a 41-dimensional sparse space, requiring L2 Ridge regularization to prevent overfitting.

15.2 System Limitations

- **Dataset Size & Static Telemetry:** The model is trained on a 7,043-record cross-sectional snapshot, which limits its ability to observe macroeconomic trends.
- **Model Precision-Recall Trade-offs:** While 80.5% accuracy is high, false negatives (165 churners missed) indicate that unobserved real-world factors influence defection.
- **Limited Behavioral Features:** Customer service tickets, network call-drop frequency, and competitor promotional offers are not present in the dataset.
- **Deployment Integration Constraints:** The web application currently operates on file-based CSV uploads and manual single inputs rather than direct live CRM API integration.

16. Conclusion

The TelcoPulse AI project achieves all core research and engineering milestones:

- **1. Successfully Developed AI/ML Solution:** An end-to-end Machine Learning customer churn decision support system was successfully conceptualized, built, and deployed.
- **2. Data Preprocessing & Analysis Completed:** Comprehensive data preprocessing, outlier auditing, and one-hot encoding were systematically implemented.
- **3. Model Rigorously Trained & Evaluated:** Multiple classifiers were trained and benchmarked, with Logistic Regression achieving 80.5% accuracy and 0.845 ROC-AUC.
- **4. Actionable Decision Support Delivered:** The system delivers real-time calibrated churn probability, annual financial exposure (\$), explainable log-odds attribution, and interactive what-if simulations.

17. Future Scope

Future enhancements planned for subsequent releases include:

- **1. Larger Datasets:** Ingest larger datasets covering multi-year customer lifecycles and millions of subscriber accounts.
- **2. Add Telemetry Features:** Integrate real-time Call Detail Records (CDR), data throughput metrics, customer service ticket sentiment, and network QoS logs.
- **3. Advanced AI/ML Ensembles:** Implement gradient boosting ensembles (XGBoost, LightGBM, CatBoost) coupled with SHAP (SHapley Additive exPlanations) values.
- **4. Real-Time Streaming Prediction:** Deploy streaming inference pipelines on customer interaction events (e.g., billing inquiries, plan downgrades).
- **5. Enterprise CRM Integration:** Direct bidirectional REST API integration with enterprise CRM suites (Salesforce, SAP Customer Experience, Zendesk).
- **6. Autonomous Retention Campaigns:** Build automated triggers that dispatch SMS or email retention offers when a subscriber enters the Critical Risk (>60%) zone.

18. References

1. IBM Corporation. 'Telco Customer Churn Dataset', IBM Business Analytics Community & Kaggle Repository, 2019.
2. Pedregosa, F., et al. 'Scikit-learn: Machine Learning in Python', Journal of Machine Learning Research, 12, pp. 2825-2830, 2011.
3. Streamlit Documentation. 'Streamlit: The fastest way to build and share data apps', <https://docs.streamlit.io/>, 2024.
4. Hastie, T., Tibshirani, R., & Friedman, J. 'The Elements of Statistical Learning: Data Mining, Inference, and Prediction', Springer, 2nd ed., 2009.
5. Ribeiro, M. T., Singh, S., & Guestrin, C. '"Why Should I Trust You?": Explaining the Predictions of Any Classifier', ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, 2016.
6. Plotly Technologies Inc. 'Collaborative data science and interactive visualization library', <https://plotly.com/python/>, 2024.
7. Hunter, J. D. 'Matplotlib: A 2D Graphics Environment', Computing in Science & Engineering, 9(3), pp. 90-95, 2007.